

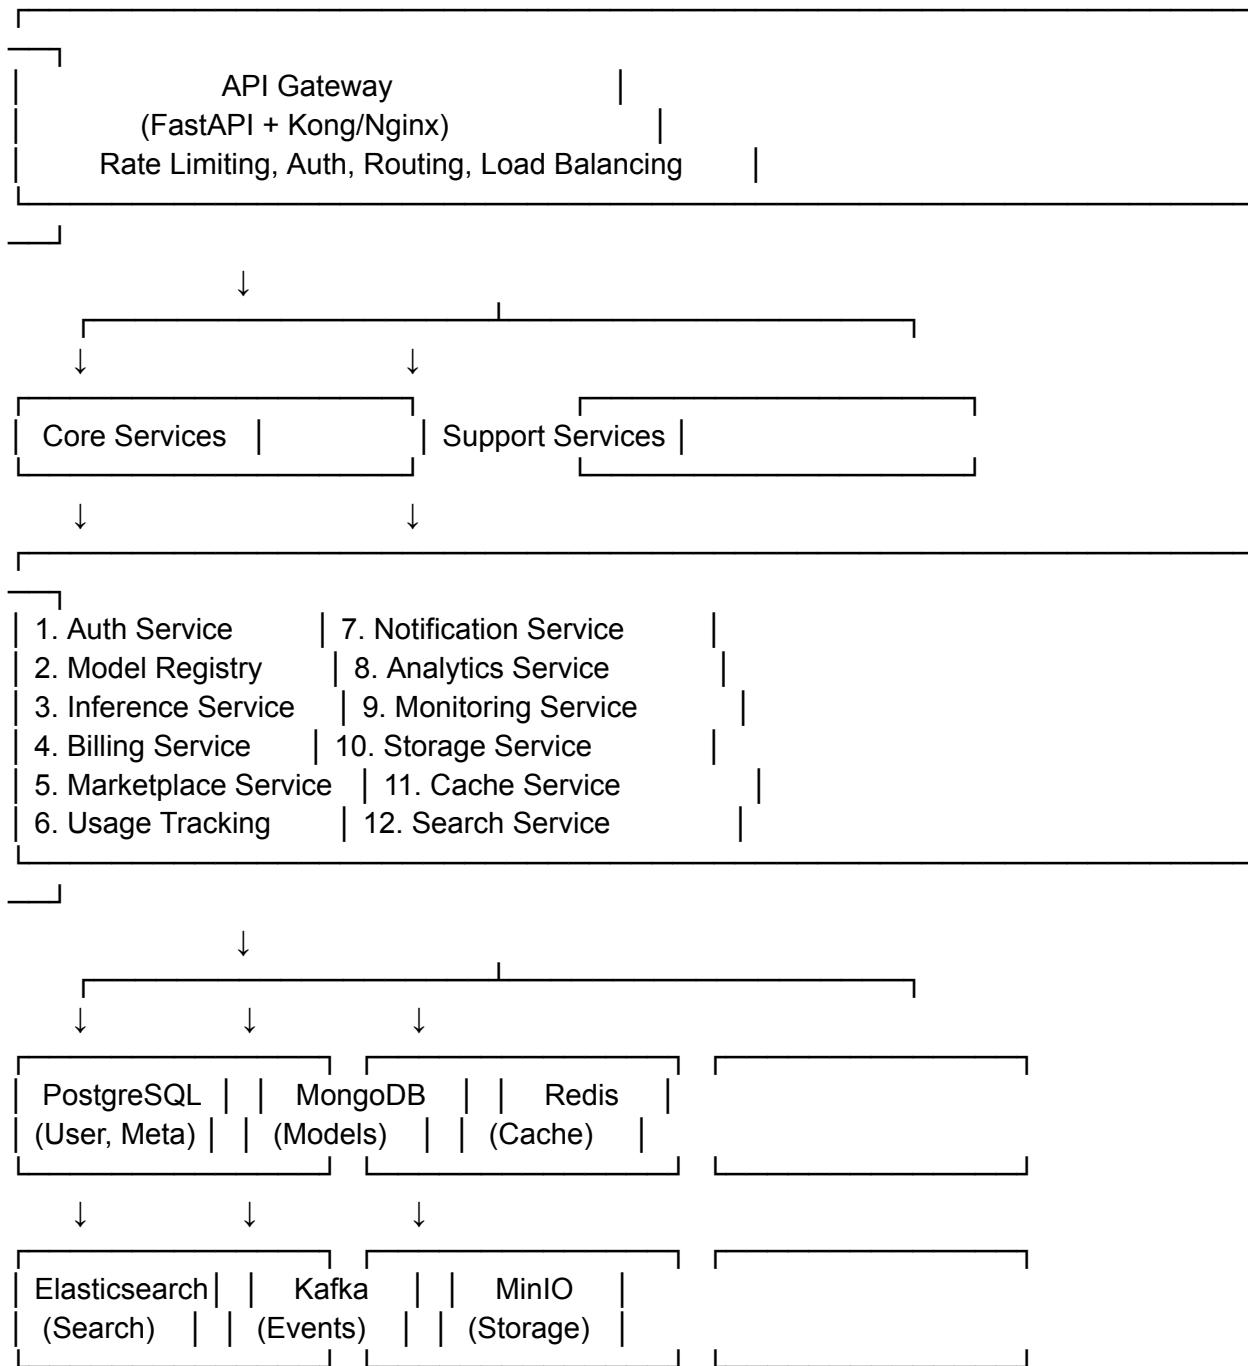
AI Model Marketplace & Inference Platform - Complete Blueprint

⌚ Project Overview

What We're Building:

A comprehensive platform where AI/ML developers can publish, monetize, and manage their models, while consumers can discover, test, and integrate AI capabilities via simple APIs - think of it as "Stripe for AI Models" or "AWS Marketplace for ML".

🏗 Complete Architecture



Complete Tech Stack Backend Framework & Languages

FastAPI (Python 3.11+) - All microservices

Pydantic - Data validation

SQLAlchemy 2.0 - ORM (async)

Alembic - Database migrations

Celery - Background task processing

AsyncIO - Async operations

Databases

PostgreSQL 15 - Primary database (users, transactions, billing)

MongoDB - Model metadata, configurations, large JSON

Redis 7 - Caching, session management, rate limiting

TimescaleDB - Time-series data (usage metrics, analytics)

Message Queue & Event Streaming

Apache Kafka - Event streaming, distributed transactions

RabbitMQ - Task queues for Celery

Redis Streams - Real-time notifications

Search & Analytics

Elasticsearch 8 - Full-text search, model discovery

Kibana - Log visualization

ClickHouse (optional) - Heavy analytics queries

Storage

MinIO - S3-compatible object storage (model files, images)

Local SSD - Hot cache for frequently used models

ML/AI Infrastructure

ONNX Runtime - Model inference (framework agnostic)

TorchServe - PyTorch model serving

TensorFlow Serving - TensorFlow models

Triton Inference Server - Multi-framework support

Ray Serve - Distributed inference (scaling)

API Gateway & Proxy

Nginx - Reverse proxy, load balancing
Kong (optional) - Advanced API gateway features
Traefik (alternative) - Cloud-native proxy

Monitoring & Observability

Prometheus - Metrics collection
Grafana - Dashboards and visualization
Jaeger - Distributed tracing
OpenTelemetry - Instrumentation
Sentry - Error tracking

Security

JWT - Authentication tokens
OAuth 2.0 - Third-party authentication
HashiCorp Vault - Secrets management
SSL/TLS - Encryption in transit
Cryptography - Data encryption at rest

Payment Processing

Stripe API - Payment gateway
Webhook handling - Payment confirmations

DevOps & Infrastructure

Docker - Containerization
Docker Compose - Local development
Kubernetes - Container orchestration (production)
Helm - Kubernetes package manager
ArgoCD - GitOps deployment
GitHub Actions - CI/CD pipelines

Testing

Pytest - Unit testing
Pytest-asyncio - Async testing
Locust - Load testing
Testcontainers - Integration testing

Documentation

Swagger/OpenAPI - API documentation
Redoc - Alternative API docs
MkDocs - User documentation

 Complete Feature Set

Phase 1: Core Platform (MVP)

1. User Management & Authentication

User registration (email/OAuth)

Email verification

Multi-factor authentication (MFA)

JWT-based authentication

API key generation and management

Role-based access control (RBAC)

Model Publishers

Model Consumers

Admins

User profiles and preferences

Password reset flows

2. Model Registry Service

Model upload (chunked upload for large files)

Model versioning (semantic versioning)

Model metadata management

Name, description, tags

Input/output schemas

Framework (PyTorch, TensorFlow, ONNX)

Model size, parameters count

License information

Model validation

Format validation

Schema validation

Security scanning

Model storage (MinIO/S3)
Model lifecycle management

Draft → Review → Published → Deprecated

3. Model Marketplace

Browse models by category

NLP (text generation, classification, NER)
Computer Vision (classification, detection, segmentation)
Audio (speech-to-text, text-to-speech)
Multimodal (CLIP, image captioning)

Search functionality (Elasticsearch)

Full-text search
Filter by tags, framework, pricing
Sort by popularity, rating, date

Model detail pages

Overview, documentation
Sample inputs/outputs
Pricing information
Usage statistics
Reviews and ratings

Model playground

Interactive testing interface
Try before you buy
Sample requests

4. Inference Service

REST API endpoints for inference

- Request validation
- Model loading and caching
- GPU/CPU resource management
- Batch inference support
- Async inference for long-running tasks
- Response formatting
- Error handling and retries

5. Billing & Subscription Service

- Pricing tiers

- Free tier (limited requests)
- Pay-as-you-go (per API call)
- Monthly subscriptions
- Enterprise plans

- Stripe integration

- Payment method management
- Subscription management
- Invoice generation

- Usage-based billing

- Track API calls per user/model
- Calculate costs dynamically

- Revenue sharing (platform fee + model owner payout)
- Billing dashboard
- Payment history

- Phase 2: Advanced Features

6. Usage Tracking & Analytics

- Real-time usage monitoring
- Request/response logging
- Performance metrics

- Latency (p50, p95, p99)

Throughput (requests/second)
Error rates

Cost tracking per model/user
Dashboard for publishers

Total requests
Revenue earned
Popular models
Geographic distribution

Dashboard for consumers

Usage statistics
Cost breakdown
Model performance

7. Rate Limiting & Quotas

Token bucket algorithm
Per-user rate limits
Per-model rate limits
Quota management

Daily/monthly limits
Soft limits (warnings)
Hard limits (rejection)

Tier-based limits
Burst allowance

8. Caching Layer

Response caching for identical requests
Cache invalidation strategies
TTL-based expiration
Cache hit/miss tracking
Distributed caching (Redis)

9. Webhook System

Webhook registration

Event types

Model published

Inference completed

Payment received

Quota exceeded

Webhook delivery

Retry mechanism with exponential backoff

Webhook verification (signatures)

10. Notification Service

Email notifications

Welcome emails

Payment confirmations

Quota warnings

Model approval/rejection

In-app notifications

Real-time alerts (WebSocket)

Notification preferences

Phase 3: Production-Grade Features

11. Model Performance Monitoring

Drift detection

Input distribution changes

Output quality degradation

A/B testing framework

Compare model versions

Traffic splitting

Statistical significance testing

Model health checks
Automated alerts for anomalies

12. Advanced Inference Features

Auto-scaling based on load
Multi-region deployment
Model ensembling
Model chaining (pipelines)
Streaming responses (for text generation)
Custom preprocessing/postprocessing
Model warm-up strategies

13. Idempotency & Reliability

Idempotency keys for requests
Request deduplication
Distributed transactions (Saga pattern)
Circuit breaker pattern
Fallback mechanisms
Graceful degradation

14. Security Features

API key rotation
IP whitelisting/blacklisting
Request signing
DDoS protection
SQL injection prevention
XSS protection
CORS configuration
Rate limit bypass protection
Audit logging

All sensitive operations logged
Immutable audit trail
Compliance reporting

15. Advanced Search & Discovery

Recommendation engine

Similar models

Personalized recommendations

Model comparison tool

Leaderboards

Most popular

Highest rated

Best performance

Collections/curated lists

16. Developer Experience

SDK generation (Python, JavaScript, Go)

Code examples

Interactive API documentation

Postman collections

CLI tool for model management

Webhooks testing tool

Sandbox environment

Phase 4: Enterprise Features

17. Multi-Tenancy

Organization accounts

Team management

Shared API keys

Organization billing

Usage aggregation

18. Custom Deployment

Private model hosting

VPC deployment

On-premise support

Dedicated infrastructure

19. Compliance & Governance

- Model provenance tracking
- Data lineage
- Compliance reports (SOC 2, GDPR)
- Data retention policies
- Right to deletion

20. Advanced Analytics

- Business intelligence dashboards
- Predictive analytics
- Cost optimization recommendations
- Anomaly detection
- Custom reports

Database Schema Design

PostgreSQL Tables

sql-- Users

```
users (
    id UUID PRIMARY KEY,
    email VARCHAR UNIQUE,
    username VARCHAR UNIQUE,
    password_hash VARCHAR,
    role VARCHAR,
    is_verified BOOLEAN,
    created_at TIMESTAMP,
    updated_at TIMESTAMP,
    deleted_at TIMESTAMP NULL -- Soft delete
)
```

-- API Keys

```
api_keys (
    id UUID PRIMARY KEY,
    user_id UUID REFERENCES users,
    key_hash VARCHAR,
    name VARCHAR,
    scopes JSON,
    last_used_at TIMESTAMP,
    expires_at TIMESTAMP,
    created_at TIMESTAMP,
    is_active BOOLEAN
)
```

```

-- Subscriptions
subscriptions (
    id UUID PRIMARY KEY,
    user_id UUID REFERENCES users,
    plan_id UUID REFERENCES plans,
    stripe_subscription_id VARCHAR,
    status VARCHAR,
    current_period_start TIMESTAMP,
    current_period_end TIMESTAMP,
    created_at TIMESTAMP
)

-- Transactions
transactions (
    id UUID PRIMARY KEY,
    user_id UUID REFERENCES users,
    amount DECIMAL,
    currency VARCHAR,
    type VARCHAR, -- charge, refund, payout
    status VARCHAR,
    stripe_payment_id VARCHAR,
    metadata JSON,
    created_at TIMESTAMP
)

-- Usage Records (TimescaleDB)
usage_records (
    id UUID,
    user_id UUID,
    model_id UUID,
    api_key_id UUID,
    request_id UUID,
    tokens_used INT,
    cost DECIMAL,
    latency_ms INT,
    status_code INT,
    timestamp TIMESTAMP,
    PRIMARY KEY (id, timestamp)
) PARTITION BY RANGE (timestamp);

-- Audit Logs
audit_logs (
    id UUID PRIMARY KEY,
    user_id UUID,

```

```
action VARCHAR,
resource_type VARCHAR,
resource_id UUID,
metadata JSON,
ip_address INET,
user_agent TEXT,
created_at TIMESTAMP
)
MongoDB Collections
javascript// Models
{
  _id: ObjectId,
  model_id: UUID,
  owner_id: UUID,
  name: String,
  description: String,
  version: String,
  framework: String, // pytorch, tensorflow, onnx
  task_type: String, // classification, generation, etc.
  tags: [String],
  // Storage
  storage_path: String,
  file_size: Number,
  checksum: String,
  // Schema
  input_schema: Object,
  output_schema: Object,
  // Configuration
  inference_config: {
    max_batch_size: Number,
    timeout_ms: Number,
    gpu_required: Boolean,
    memory_mb: Number
  },
  // Pricing
  pricing: {
    free_tier_limit: Number,
    price_per_request: Number,
    price_per_token: Number,
    subscription_prices: Object
```

```
    },  
  
    // Metrics  
    stats: {  
        total_requests: Number,  
        avg_latency_ms: Number,  
        success_rate: Number,  
        rating: Number,  
        review_count: Number  
    },  
  
    // Lifecycle  
    status: String, // draft, review, published, deprecated  
    published_at: Date,  
    deprecated_at: Date,  
  
    created_at: Date,  
    updated_at: Date  
}  
  
// Model Reviews  
{  
    _id: ObjectId,  
    model_id: UUID,  
    user_id: UUID,  
    rating: Number,  
    comment: String,  
    created_at: Date  
}  
  
// Inference Cache  
{  
    _id: ObjectId,  
    model_id: UUID,  
    request_hash: String, // Hash of input  
    response: Object,  
    ttl: Date,  
    created_at: Date  
}  
...  
  
---
```

1. API Gateway Service

...

Port: 8000

Technology: FastAPI + Nginx

Responsibilities:

- Request routing
- Authentication validation
- Rate limiting
- Request/response logging
- CORS handling
- SSL termination

...

2. Auth Service

...

Port: 8001

Technology: FastAPI + PostgreSQL + Redis

Responsibilities:

- User registration/login
- JWT token generation
- API key management
- OAuth 2.0 flows
- MFA handling
- Session management

...

3. Model Registry Service

...

Port: 8002

Technology: FastAPI + MongoDB + MinIO

Responsibilities:

- Model CRUD operations
- Version management
- File upload handling
- Model validation
- Metadata management
- Search indexing

...

4. Inference Service

...

Port: 8003

Technology: FastAPI + ONNX/TorchServe + Redis

Responsibilities:

- Load models into memory
 - Process inference requests
 - Batch processing
 - Result caching
 - Resource management
 - Performance optimization
- ...

5. Billing Service

...

Port: 8004

Technology: FastAPI + PostgreSQL + Stripe

Responsibilities:

- Subscription management
 - Payment processing
 - Invoice generation
 - Usage calculation
 - Revenue distribution
 - Webhook handling
- ...

6. Usage Tracking Service

...

Port: 8005

Technology: FastAPI + TimescaleDB + Kafka

Responsibilities:

- Log all API requests
 - Calculate usage metrics
 - Generate analytics
 - Quota management
 - Cost tracking
- ...

7. Marketplace Service

...

Port: 8006

Technology: FastAPI + MongoDB + Elasticsearch

Responsibilities:

- Model discovery
- Search and filtering
- Recommendations
- Reviews and ratings
- Model playground

...

8. Notification Service

...

Port: 8007

Technology: FastAPI + RabbitMQ + Redis

Responsibilities:

- Email sending
- Webhook delivery
- Push notifications
- Alert management
- Notification preferences

...

9. Analytics Service

...

Port: 8008

Technology: FastAPI + TimescaleDB + ClickHouse

Responsibilities:

- Dashboard data
- Report generation
- Trend analysis
- Anomaly detection
- Business intelligence

...

10. Storage Service

...

Port: 8009

Technology: FastAPI + MinIO

Responsibilities:

- File upload/download
- Presigned URL generation
- File versioning
- Storage optimization

Development Roadmap

Week 1-2: Project Setup & Foundation

Set up project structure

Docker Compose for local development

Database setup (PostgreSQL, MongoDB, Redis)

Basic FastAPI gateway

Authentication service (JWT)

User CRUD operations

Week 3-4: Model Registry

Model upload endpoint (chunked)

MinIO integration

MongoDB schema for models

Model versioning

Basic validation

Week 5-6: Inference Service

ONNX Runtime integration

Model loading mechanism

Basic inference endpoint

Response caching (Redis)

Error handling

Week 7-8: Billing Integration

Stripe integration

Subscription plans

Usage tracking basics

Payment webhooks

Week 9-10: Marketplace

Model listing API

Search with Elasticsearch

Model detail pages

Playground interface

Week 11-12: Advanced Features

Rate limiting

Idempotency

Circuit breaker

Distributed tracing

Monitoring setup

Week 13-14: Event-Driven Architecture

Kafka setup

Event producers/consumers

Saga pattern for transactions
Webhook system

Week 15-16: Testing & Optimization

Unit tests
Integration tests
Load testing
Performance optimization
Documentation

Week 17-18: DevOps & Deployment

Kubernetes manifests
CI/CD pipelines
Monitoring dashboards
Production deployment

Key Metrics to Track Business Metrics

Total users (publishers + consumers)
Active models
Total API requests
Revenue (MRR, ARR)
Churn rate
Customer acquisition cost

Technical Metrics

API latency (p50, p95, p99)
Error rate
Uptime (SLA: 99.9%)
Cache hit rate
Database query performance
Message queue lag

User Metrics

Model upload success rate
Average inference time
Cost per request
Most popular models

User retention

⌚ Success Criteria

After completing this project, you'll have:

- Portfolio Project - Production-grade system for interviews
- Microservices Expertise - Built 10+ interconnected services
- Event-Driven Mastery - Kafka, Saga pattern, CQRS
- Payment Integration - Real Stripe implementation
- Scalability Knowledge - Caching, sharding, load balancing
- DevOps Skills - Docker, Kubernetes, CI/CD
- ML Integration - AI model serving at scale
- Advanced Patterns - Circuit breaker, idempotency, distributed tracing